

Fast and simple connectivity in graph timelines

Adam Karczmarz ^{*} and Jakub Łącki [†]

University of Warsaw

{a.karczmarz|j.lacki}@mimuw.edu.pl

Abstract

In this paper we study the problem of answering connectivity queries about a *graph timeline*. A graph timeline is a sequence of undirected graphs G_1, \dots, G_t on a common set of vertices of size n such that each graph is obtained from the previous one by an addition or a deletion of a single edge. We present data structures, which preprocess the timeline and can answer the following queries:

- **forall**(u, v, a, b) – does the path $u \rightarrow v$ exist in *each* of G_a, \dots, G_b ?
- **exists**(u, v, a, b) – does the path $u \rightarrow v$ exist in *any* of G_a, \dots, G_b ?
- **forall2**(u, v, a, b) – do there exist two edge-disjoint paths connecting u and v in *each* of G_a, \dots, G_b ?

We show data structures that can answer **forall** and **forall2** queries in $O(\log n)$ time after preprocessing in $O(m + t \log n)$ time. Here by m we denote the number of edges that remain unchanged in each graph of the timeline. For the case of **exists** queries, we show how to extend an existing data structure to obtain a preprocessing/query trade-off of $\langle O(m + \min(nt, t^{2-\alpha})), O(t^\alpha) \rangle$ and show a matching conditional lower bound.

1 Introduction

In this paper we revisit the problem of maintaining the connectivity information in a *graph timeline*. The problem was formulated and solved **in a recent paper by Łącki and Sankowski [9]**. They define a graph timeline to be a sequence of graphs G_1, G_2, \dots, G_t on a common set of vertices V of size n such that the graph G_i is obtained from G_{i-1} by adding or deleting a single edge. Their goal was to preprocess the graph timeline to build a data structure that may answer connectivity queries regarding a contiguous fragment of the timeline:

- **forall**(u, v, a, b) — are vertices u and v connected by a path in *each* of G_a, G_{a+1}, \dots, G_b ?
- **exists**(u, v, a, b) — are vertices u and v connected by a path in *any* of G_a, G_{a+1}, \dots, G_b ?

We stress that the entire timeline is revealed in the very beginning for preprocessing, and after that the queries may arrive in an online fashion.

Throughout this paper, we write $\langle f(n, m, t), g(n, m, t) \rangle$ to denote a data structure, whose **preprocessing time is $f(n, m, t)$ and the query time is $g(n, m, t)$** .

^{*}Supported by the grant NCN2014/13/B/ST6/01811 of the Polish Science Center. Partially supported by FET IP project MULTIPLEX 317532.

[†]Jakub Łącki is a recipient of the Google Europe Fellowship in Graph Algorithms, and this research is supported in part by this Google Fellowship.

In the case of **forall** queries, Łącki and Sankowski presented an $\langle O(m + t \log t \log \log t \log n), O(\log n \log \log t) \rangle$ data structure. Here by m we denote the number of edges that remain unchanged in each of G_1, \dots, G_t . Their data structure is Monte Carlo randomized and the query time is amortized. For **exists** queries they give an $\langle O(m + nt), O(1) \rangle$ data structure.

We improve the results of [9] and show new algorithms, which are more efficient, simpler and deterministic. In addition, we also develop an extended data structure that may efficiently answer an even more complex query regarding 2-edge-connectivity:

- **forall2**(u, w, a, b) — are vertices u and v connected by two edge-disjoint paths in *each* of G_a, G_{a+1}, \dots, G_b ?

Moreover, we give new conditional lower bounds for the problem of answering **exists** queries, which also improves the results of [9].

1.1 Related work

A rich body of connectivity-related dynamic problems has been studied in the area of networks and distributed computing. A number of such problems has been surveyed in [2]. In a typical scenario, we work with a sequence of graphs $G^t = G_1, \dots, G_t$ that represent the states of an evolving network at different points in time. However, the properties of these graphs, which are of interest, such as *T-interval connectivity* [8] or *time-respecting paths* [7] are usually much more complex than what can be studied with ordinary connectivity queries, that is queries about the existence of a path connecting two given vertices in a particular graph. For example, the problem of T-interval connectivity consists of deciding if for every subsequence G_a, \dots, G_{a+T-1} of T consecutive graphs in G^t , the intersection $G_a \cap \dots \cap G_{a+T-1}$ of these graphs contains a connected component spanning all vertices. Here we define the intersection of two graphs to be the graph obtained by intersecting their edge sets.

We believe that the queries we consider in this paper are powerful enough to study interesting properties of evolving networks. A **forall** query checks if two vertices are connected with a path in every graph among G_a, \dots, G_b , but the path can be different in each of the graphs and may not even exist in the intersection of these graphs. Even stronger is a **forall2** query, checking whether two vertices are connected with two edge-disjoint paths in each graph of the given fragment. This may serve as a measure of robustness of connection between two nodes of a network.

The algorithms that process graph timelines can also be considered *semi-offline* counterparts of dynamic graph algorithms. The updates are given upfront, but the queries may arrive in an online fashion, i.e. they are issued one by one, only after the preprocessing is finished. A possible scenario for the semi-offline model would be to collect and index the history of evolving network up to some point of time and then use the queries to analyze various properties of the network efficiently.

It is worth noting that the knowledge of the entire history of changes in most cases leads to data structures faster and simpler than the best online ones. However, this property has rarely been exploited to design efficient algorithms. Eppstein [4] has shown an algorithm, which, given a weighted graph G and a sequence of k edge weight updates, computes the weight of the minimum spanning tree after each update in $O((m + k) \log n)$ time.

1.2 Our results

We show $\langle O(m + t \log n), O(\log n) \rangle$ data structures for answering **forall** queries and **forall2** queries. The data structures use $O(t \log n)$ space. This improves the results of [9] in a number of ways: our algorithms are faster and deterministic, use less space, the time bounds are worst-case

and the query time is independent of the length of the timeline. We also introduce **forall2** queries, which were not considered before. On top of that, our algorithms are arguably simpler.

What is interesting, we obtain a solution for the 2-edge-connectivity problem, which is much more efficient than what has been achieved in the dynamic case. The best known algorithm for 2-edge-connectivity is due to Holm et al. [5]. It processes t updates in $O((t + m) \log^4 n)$ time, where m is the initial number of edges, and answers queries in $O(\log n)$ time. Our algorithm may preprocess the timeline in only $O(m + t \log n)$ time to answer queries in $O(\log n)$ time.

In the construction of the algorithm for answering **forall** queries we use the following two observations. Consider a timeline G_1, \dots, G_t . If there is an edge uw present in every graph among G_1, \dots, G_t , vertices u and w are equivalent from the point of view of any query, so the edge uw can be contracted in each graph. Once we do that, we are left with $O(t)$ edges in total, each being added or deleted at some point of time. Thus, if there are much more than t vertices, some vertices are isolated in every G_1, \dots, G_t , and can be safely treated separately in the beginning and removed. These ideas are then used recursively in a divide-and-conquer algorithm, which at each step halves the length of the timeline to compute a segment tree over the sequence G_1, \dots, G_t . This segment tree stores connectivity information about every individual graph in the timeline. Here we adapt the ideas of Eppstein's reduction and contraction scheme used for offline computation of minimum spanning trees [4].

Next, we use a fingerprinting scheme to identify vertices belonging to the same connected components in multiple consecutive graphs, which allows us to answer **forall** queries. Additionally, our fast algorithm for answering queries uses a data structure for efficient testing of equality of contiguous subsequences of a given sequence. This is then extended to handle **forall2** queries.

For **exists** queries, we show how to leverage the $\langle O(m + nt), O(1) \rangle$ data structure from [9] to build an $\langle O(m + \min(nt, t^{2-\alpha})), O(t^\alpha) \rangle$ data structure, where α is a parameter from the range $[0, 1)$, which can be chosen arbitrarily. All of the presented algorithms are simple and can easily be implemented.

Moreover, we develop a conditional lower bound for the problem of answering **exists** queries. We show that answering t **exists** queries on a timeline of length t , consisting of graphs with $O(t)$ edges, can be used to detect triangles in a graph with $O(t)$ edges. This implies a conditional lower bound of $\Omega(t^{1.41})$ and improves the result of [9], where a weaker lower bound was shown. We also show that an $O(t^{1.5-\epsilon})$ *combinatorial* algorithm for the aforementioned problem would imply a subcubic *combinatorial* algorithm for the Boolean matrix multiplication problem, which would be a major breakthrough. At the same time, our improved data structure for **exists** queries may solve this problem in $O(t^{1.5})$ time, which means that it is, in some sense, optimal.

1.3 Organization of this paper

In Section 2 we introduce notation and give a few simple properties of segment trees, which we later use. Section 3 describes the basic version of our data structure, which is then extended to handle **forall** and **forall2** queries. Then, in Section 4 we present an algorithm for answering **forall** queries. Next, in Section 5 we develop improved lower bounds for the problem of answering **exists** queries, as well as show that a trade-off between query and preprocessing time is possible. Finally, in Section 6 we discuss the possible directions of future research.

2 Preliminaries

A *graph timeline* is a sequence G^t of graphs G_1, G_2, \dots, G_t , where $G_i = (V, E_i)$. We call each individual graph in G^t a *version*. For each $i \in [1, t)$ we have $|E_i \oplus E_{i+1}| = 1$, i.e. E_{i+1} is

obtained from E_i by adding or deleting a single edge. We assume that the input is given as the set E_1 and a list of $t - 1$ operations that describe, for each $i \in [1, t - 1]$, how to obtain E_{i+1} from E_i .

Throughout this paper we work with intervals of integers, that is $[a, b]$ denotes $\{a, a + 1, \dots, b\}$. We say that edge (u, v) is *alive* in the interval $[x, y]$ iff $(u, v) \in E_j$ for each $j \in [x, y]$. For each edge $e \in E_1 \cup \dots \cup E_t$ we define $L(e)$ to be the set of maximal intervals such that e is alive in each of them. An edge e is called *permanent* iff $L(e) = \{[1, t]\}$, that is, it is present in every version. Otherwise, we say that e is a *temporary* edge. We denote by m the number of permanent edges. The number of temporary edges is at most t . We begin the initialization of our data structures by finding the sets $L(e)$ in $O(|E_1| + t) = O(m + t)$ time.

We denote by Δ_a^+ the set of edges e such that $[a, x] \in L(e)$ for some $x \in [a, t]$, i.e., edges present in G_a , but not in G_{a-1} . Similarly, let Δ_b^- be the set of edges e such that $[x, b] \in L(e)$ for some $x \in [1, b]$. It is easy to verify that $\sum_{i=1}^t |\Delta_i^+| + \sum_{i=1}^t |\Delta_i^-| = O(m + t)$. Moreover, for $a \in (1, t]$, we have $|\Delta_a^+| \leq 1$, while for $b \in [1, t)$ we have $|\Delta_b^-| \leq 1$.

Throughout the paper, we assume that $t \geq n$ and $t = 2^B$ for some integer $B \geq 0$. The latter assumption can be achieved by adding dummy graphs to the timeline.

2.1 Elementary intervals and the segment tree

Given $t = 2^B$, the set of *elementary intervals* is defined inductively:

1. $[1, t]$ is an elementary interval,
2. if $[a, b]$ is an elementary interval, and $a < b$ we let $\text{mid} = \lfloor \frac{a+b}{2} \rfloor$, and define $[a, \text{mid}]$ and $[\text{mid} + 1, b]$ to be elementary intervals as well.

The set of elementary intervals can be naturally organized into a complete binary tree, which we call a *segment tree*. Assuming the above notation, we call $\text{left}([a, b]) = [a, \text{mid}]$ the left child of interval $[a, b]$. Similarly, $\text{right}([a, b]) = [\text{mid} + 1, b]$. The parent interval of P is denoted by $\text{par}(P)$. We first prove a few properties of elementary intervals.

Lemma 1. *Every two elementary intervals are either disjoint, or one of them is contained in the other. The latter is the case iff one of them is a descendant of the other in the segment tree.*

Lemma 2. *Every interval $[c, d] \subseteq [1, t]$ can be partitioned into no more than $2 \log_2(d - c + 1) + 2$ disjoint elementary intervals such that no two intervals from the partition can be merged into a bigger elementary interval. The partition can be computed in time $O(\log(d - c + 1))$.*

Proof. If $c = d$, then the interval does not have to be partitioned at all. Assume $c < d$. Consider the leaves $[c, c]$ and $[d, d]$ of the segment tree and let P be the lowest common ancestor of these intervals, i.e., the smallest elementary interval which contains both c and d . Our initial partition is formed by the following intervals:

- $[c, c]$ and $[d, d]$,
- if both the interval Q and its parent lie on the path from $[c, c]$ to P (but excluding P) and also $Q = \text{left}(\text{par}(Q))$, we include $\text{right}(\text{par}(Q))$ (i.e. the sibling of Q) in our partition,
- if both the interval Q and its parent lie on the path from $[d, d]$ to P (but excluding P) and also $Q = \text{right}(\text{par}(Q))$, we include $\text{left}(\text{par}(Q))$ (i.e. the sibling of Q) in our partition.

We first show that the chosen family of intervals W is indeed a partition of $[c, d]$. By Lemma 1, the chosen intervals are disjoint, since there are no two such that one of them is an ancestor

of the other. For any interval from W , its left endpoint is not less than c , whereas its right endpoint is not larger than d . Hence, $\bigcup W \subseteq [c, d]$. Moreover, $[c, d] \subseteq \bigcup W$. It is clear that $\{c, d\} \subseteq \bigcup W$. To show that $f \in (c, d)$ belongs to $\bigcup W$ consider a path from $[f, f]$ to P . This path either joins the path $[c, c] \rightarrow P$ from the right, or joins the path $[d, d] \rightarrow P$ from the left. In both of these cases, the last interval Q of $[f, f] \rightarrow P$ before the paths merged ($f \in Q \subseteq [c, d]$) was included in W .

Let us count the number of intervals in W . First notice, that every elementary interval in W is not longer than $d - c + 1$. Furthermore, each subsequent interval chosen from one of the paths ($[c, c] \rightarrow P$ or $[d, d] \rightarrow P$) is at least twice as long as the previous interval taken while climbing that path. Taking into account the additional intervals $[c, c]$ and $[d, d]$, we get the bound $2 \log_2(d - c + 1) + 2$. The $O(\log(d - c + 1))$ time can be achieved by climbing the two paths simultaneously.

The above procedure does not guarantee that no two elementary intervals from W can be merged into a larger elementary interval. However, this can be easily fixed. Every time when we put into W an elementary interval such that its sibling in the segment tree is already contained in W , we replace the two siblings with their parent. As the lengths of the elementary intervals put into W only increase on a path $[c, c] \rightarrow P$ or $[d, d] \rightarrow P$, the potential sibling can only be the interval that was the last to be included in W .

Eventually, we might also end up with $W = \{\text{left}(P), \text{right}(P)\}$; then we ought to replace the partition with $\{P\}$.

This fix does not influence the overall time complexity of the partitioning, which remains $O(\log(d - c + 1))$. \square

Lemma 3. *If P_1, P_2, \dots, P_k are disjoint intervals contained in $[1, t]$, we can partition them into at most $2k (\log_2 \frac{t}{k} + 1)$ disjoint elementary intervals.*

Proof. We use Lemma 2 to partition each of P_1, \dots, P_k . For $1 \leq i \leq k$, let $l_i = |P_i|$. Since the intervals are disjoint, their partitions into elementary intervals are also disjoint. Hence, by Lemma 2, the total size of the partition can be bounded as follows:

$$2k + 2 \sum_{i=1}^k \log_2 l_i \leq 2k + 2k \log_2 \left(\frac{1}{k} \sum_{i=1}^k l_i \right) \leq 2k + 2k \log_2 \frac{t}{k} = 2k \left(\log_2 \frac{t}{k} + 1 \right).$$

We used the bound $\sum_{i=1}^k l_i \leq t$ and the Jensen's inequality for the concave function $f(x) = \log_2 x$. \square

As it is much easier to work with elementary intervals, for each edge e we partition all intervals from $L(e)$ into elementary intervals.

Lemma 4. *All intervals in $\bigcup_{e \in V \times V} L(e)$ can be partitioned into $O(m + t \log n)$ elementary intervals. The partition can be performed in time $O(m + t \log n)$.*

Proof. Denote by E^* the set of temporary edges. For any $e \in E^*$, let us denote by q_e the number $|L(e)|$. We have $\sum_{e \in E^*} q_e \in [\frac{t}{2}, t]$ and $|E^*| \leq \min(t, n^2)$. By Lemma 3, we conclude



that the total number of elementary intervals for temporary edges is at most

$$\begin{aligned}
2 \sum_{e \in E^*} \left(q_e \log_2 \left(\frac{t}{q_e} \right) + 1 \right) &\leq 2|E^*| \left(\sum_{e \in E^*} \frac{q_e}{|E^*|} \log_2 \left(\frac{t}{q_e} \right) \right) + 2t \\
&\leq 2|E^*| \left(\frac{1}{|E^*|} \sum_{e \in E^*} q_e \right) \log_2 \left(\frac{t}{\left(\frac{1}{|E^*|} \sum_{e \in E^*} q_e \right)} \right) + 2t \\
&\leq 2t \log_2 (2|E^*|) + 2t \\
&= O(t \log n).
\end{aligned}$$

Here we used the Jensen's inequality for the concave function $f(x) = x \log_2 \frac{t}{x}$ and weights equal to $\frac{1}{|E^*|}$. Since each permanent edge has exactly one interval in its partition, we obtain the desired bound $O(m + t \log n)$. \square

For an elementary interval $[a, b]$, we set $E_{[a, b]}$ to be the set of edges that contain $[a, b]$ in their partition. From Lemmas 2 and 4 it follows that each edge is contained in $O(\log t)$ sets $E_{[a, b]}$ and the sum over elementary intervals $\sum_{[a, b]} E_{[a, b]}$ is of order $O(m + t \log n)$.

3 The data structure

We now describe a tree-like data structure T , which is a crucial part of all our algorithms. In the following we reserve the name T for this particular data structure. The data structure T is based on the set of all elementary intervals organized into a complete binary tree. This tree has a single node $T_{[a, b]}$ for each elementary interval $[a, b]$. Denote by $G_{[a, b]}$ the graph $(V, E_a \cap \dots \cap E_b)$. Roughly speaking, our goal is to associate with $T_{[a, b]}$ the information about the connected components of $G_{[a, b]}$. We first give a simple approach for constructing the data structure T , and then show how to speed it up. We use the following fact.

Lemma 5. *Let $[a, b]$ be an elementary interval such that $[a, b] \neq [1, t]$. Then $E(G_{[a, b]}) = E(G_{\text{par}([a, b])}) \cup E_{[a, b]}$.*

Proof. Recall that $E(G_{[a, b]}) = E(G_a) \cap \dots \cap E(G_b)$. Thus, $E(G_{\text{par}([a, b])}) \subseteq E(G_{[a, b]})$. Moreover, directly from the definitions we have $E_{[a, b]} \subseteq E(G_{[a, b]})$. It remains to show $E(G_{[a, b]}) \subseteq E(G_{\text{par}([a, b])}) \cup E_{[a, b]}$.

Consider an edge $e \in E(G_{[a, b]})$. If e is alive in some interval $[c, d] \supseteq [a, b]$ such that $\text{par}([a, b]) \subseteq [c, d]$, we have $e \in E(G_{\text{par}([a, b])})$. To complete the proof it remains to consider the case when e is alive in an interval $[c, d]$, and $\text{par}([a, b]) \not\subseteq [c, d]$. We show that in this case $e \in E_{[a, b]}$. In other words, $[a, b] \in X$, where X is the partition of elements of $L(e)$ into elementary intervals.

From $\text{par}([a, b]) \not\subseteq [c, d]$ it follows that no ancestor of $[a, b]$ is a part of X . At the same time, every elementary interval that is neither an ancestor nor a descendant of $[a, b]$ is disjoint with $[a, b]$ (Lemma 1), so it does not belong to X either. Therefore, each elementary interval that belongs to X and intersects $[a, b]$ is either $[a, b]$ or one of its descendants. Consequently, X contains a subset of disjoint elementary intervals Y such that $\bigcup Y = [a, b]$. Suppose $Y \neq \{[a, b]\}$. Let Q_1 be the shortest interval from Y . Note that $Q_1 \neq [a, b]$. Also, by Lemma 2, the sibling Q_2 of Q_1 is not contained in Y . As Q_1 is the shortest, no descendant of Q_2 is contained in Y . Moreover, as $Q_1 \in Y$, no ancestor of Q_2 is contained in Y . Hence, Y does not cover the integers from Q_2 , a contradiction. Thus, we have $Y = \{[a, b]\}$, so $[a, b] \in X$ and, as a result, we have $e \in E_{[a, b]}$. \square

In the simple approach, we associate with $T_{[a,b]}$ a graph $S_{[a,b]}$, which has a single vertex for each connected component of $G_{[a,b]}$, and does not contain any edges. By Lemma 5, $G_{[a,b]}$ is obtained from $G_{\text{par}([a,b])}$ by adding some edges. This implies that each component of $G_{[a,b]}$ is a sum of some components of $G_{\text{par}([a,b])}$. To compute $S_{[a,b]}$ we build a graph H on a vertex set $V(S_{\text{par}([a,b])})$ and add to it edges of $E_{[a,b]}$ (each edge endpoint has to be mapped to its connected component in $G_{\text{par}([a,b])}$) and then find its connected components. These components are exactly the components of $G_{[a,b]}$. Observe that during this computation we may also compute a mapping between the vertices of $S_{\text{par}([a,b])}$ and $S_{[a,b]}$. In the case of $S_{[1,t]}$ we compute a mapping between individual vertices and connected components of $G_{[1,t]}$.

T represents the connected components of every graph in the timeline. Consider a graph G_c . In order to find a connected component of a vertex v in G_c , we traverse the path in T from $T_{[1,t]}$ to $T_{[c,c]}$. We compute the connected component of vertex v in every graph $G_{[a,b]}$ on the path. Observe that if we know the connected component of v in $G_{\text{par}([a,b])}$, we may compute the connected component of v in $G_{[a,b]}$ by following the mapping between the components of $G_{\text{par}([a,b])}$ and $G_{[a,b]}$. At the end of the traversal, we find the component of v in $G_{[c,c]} = G_c$.

3.1 An efficient construction

In order to compute the data structure T efficiently, we need to make an additional optimization, which is crucial for obtaining good running time.

Consider an elementary interval $[a, b]$ and a connected component C of $G_{[a,b]}$. Assume that within the graphs G_a, \dots, G_b no edge incident to a vertex of C is ever added or deleted. In other words, the edges incident to vertices of C are the same in each of G_a, \dots, G_b . This means that in each of G_a, \dots, G_b vertices of C are connected to each other, but not connected to *any* vertex outside C . Hence, C is also a connected component in each of G_a, \dots, G_b .

As a result, there is no need to store C in the descendants of $T_{[a,b]}$. When searching for a connected component of a vertex $v \in C$ in G_c , where $c \in [a, b]$, we may simply stop the search in the representation of C in $T_{[a,b]}$. This observation will be used in the reduction phase of the construction of the tree T .

We now describe the efficient construction of the tree T . For each node $T_{[a,b]}$ of T , where $[a, b]$ is an elementary interval, we compute a graph $S_{[a,b]}$. The vertices of $S_{[a,b]}$ correspond to *some* of the components of $G_{[a,b]}$. We say that $v \in V$ is *represented* in $S_{[a,b]}$ if there is a vertex $s \in V(S_{[a,b]})$ that corresponds to a component containing v . The graphs $S_{[a,b]}$ have no edges.¹

Let $[a, b]$ be an elementary interval. $S_{[a,b]}$ is computed based on $S_{\text{par}([a,b])}$ (or (V, \emptyset) , if $[a, b] = [1, t]$) in two phases called *reduction* and *contraction*.

In the reduction phase some vertices of $H = S_{\text{par}([a,b])}$ are removed, as they are not affected by any edge addition or deletion that is carried out among G_a, \dots, G_b . Namely, we mark endpoints of edges in $F = E_{[a,b]} \cup \bigcup_{i=a+1}^b \Delta_i^+ \cup \bigcup_{i=a}^{b-1} \Delta_i^-$ and then remove the unmarked vertices. Note that the sets $E_{[a,b]}$, Δ_i^+ and Δ_i^- contain edges of the original graph, so their endpoints have to be mapped to the corresponding vertices of H . The reduction phase is performed only when $b - a + 1 < n$. It is done by a call $\text{REDUCE}(H, F)$, which produces a pair (S', M) , where S' is the reduced graph and M is a mapping between $V(S_{\text{par}([a,b])})$ and $V(S') \cup \{\perp\}$. The value of \perp means that a vertex has been removed and does not have a corresponding vertex in S' . The procedure can be implemented with a simple graph search to work in $O(|H| + |F|)$ time.

In the second phase, called the contraction phase, some of the remaining vertices of $H = S'$ are merged to form $S_{[a,b]}$. Specifically, the components formed in S' after adding edges $F = E_{[a,b]}$ are contracted. Again, we use a function $\text{CONTRACT}(H, F)$, which produces a pair (S', M)

¹Defining a graph with no edges may look confusing. However, we define $S_{[a,b]}$ to be a graph, as we add edges to $S_{[a,b]}$ in our data structure for 2-edge-connectivity.

consisting of the contracted graph S' and the mapping between H and S' . This function can also be easily implemented to work in linear time.

Consider an elementary interval P . Together with S_P , the node T_P stores two tables \mathbf{l}_P and \mathbf{r}_P mapping vertices of S_P to $V(S_{\text{left}(P)}) \cup \{\perp\}$ and $V(S_{\text{right}(P)}) \cup \{\perp\}$ respectively. If $\mathbf{l}_P[k] \neq \perp$, $\mathbf{l}_P[k]$ is the vertex of $S_{\text{left}(P)}$ that corresponds to $k \in V(S_P)$. $\mathbf{l}_P[k] = \perp$ means that P is a leaf, or there is no vertex corresponding to k in $S_{\text{left}(P)}$. The table \mathbf{r}_P is defined analogously. For simplicity, we also assume that $T_{[1,t]}$ is a left child of a special node $T_{[0,\infty]}$ and $S_{[0,\infty]} = (V, \emptyset)$, so that for each $v \in V$, $\mathbf{l}_{[0,\infty]}[v]$ points to the vertex of $S_{[1,t]}$ representing the original vertex v .

The graphs S_P along with \mathbf{l} and \mathbf{r} pointers are sufficient to find the component of any vertex v in any of G_1, \dots, G_t . To access the component of vertex v in G_c we start at vertex v in $S_{[0,\infty]}$ and follow \mathbf{l} or \mathbf{r} pointers in order to reach the leaf $T_{[c,c]}$. The traversal stops once we reach $T_{[c,c]}$ or the pointer we want to use ($\mathbf{l}[k]$ or $\mathbf{r}[k]$) is equal to \perp . Let P be the elementary interval, where the traversal finishes and k be the vertex in S_P , which we reached. Then, as we later show, (k, P) uniquely identifies the component of vertex v in G_c . The above process can be seen as a function $\text{COMP-ID}(w, a, b, c)$ that follows the path to $T_{[c,c]}$ starting at vertex $w \in V(S_{[a,b]})$. The pair (k, P) , defined as above, is what the call $\text{COMP-ID}(\mathbf{l}_{[0,\infty]}[v], 1, t, c)$ returns. The full text of the COMP-ID function is given in Appendix A.

Lemma 6. *Let $1 \leq c \leq t$. For any $u \in V$, denote by (k_u, P_u) the value returned by $\text{COMP-ID}(\mathbf{l}_{[0,\infty]}[u], 1, t, c)$. Then, two vertices $v, w \in V$ are connected by a path in G_c iff $k_v = k_w$ and $P_v = P_w$.*

Proof. Observe that $[c, c] \subseteq P_v$ and $[c, c] \subseteq P_w$. First, assume that $P_v \neq P_w$, which, by Lemma 1 means that one of the intervals contains the other one. Without loss of generality suppose that $P_v \subseteq \text{left}(P_w)$. Then k_w is a component of G_{P_w} that is not incident to any changes in the time interval $\text{left}(P_w)$, while v is in some component of G_{P_w} that undergoes changes in $\text{left}(P_w)$. Thus, these are different components. If $P_v = P_w$, then both k_v and k_w are components of G_{P_v} not incident to any changes in the time interval P_v . Both v and w , however, are represented in S_{P_v} , so they are in the same component iff $k_v = k_w$. \square

Let us bound the time needed to build T . We begin with an auxiliary lemma, whose proof is based on the fact that we perform the reduction.

Lemma 7. *Let $[a, b]$ be an elementary interval. Then $|V(S_{[a,b]})| \leq \min(8(b - a + 1), n)$.*

Proof. Clearly, $|V(S_{[a,b]})| \leq n$. For $b - a + 1 \geq n$, we have $8(b - a + 1) \geq n$, so $|V(S_{[a,b]})| \leq n \leq \min(n, 8(b - a + 1))$ holds. Assume that $b - a + 1 < n$. The graph $S_{[a,b]}$ is constructed from $S_{\text{par}([a,b])}$ by applying reduction and contraction. Let $C = \bigcup_{i=a+1}^b \Delta_i^+ \cup \bigcup_{i=a}^{b-1} \Delta_i^-$. The reduction produces a graph S' of at most $2|E_{[a,b]} \cup C|$ vertices. The contraction does not increase the number of vertices. Therefore, $|V(S_{[a,b]})| \leq 2|E_{[a,b]} \cup C|$.

Let $\text{par}([a, b]) = [a_1, b_1]$ and consider the analogous set C_1 for $\text{par}([a, b])$, i.e. $C_1 = \bigcup_{i=a_1+1}^{b_1} \Delta_i^+ \cup \bigcup_{i=a_1}^{b_1-1} \Delta_i^-$. For $1 < i \leq t$ we have $|\Delta_i^+| \leq 1$ and for $1 \leq i < t$ we have $|\Delta_i^-| \leq 1$. Moreover, since $\text{par}([a, b]) = [a_1, b_1]$, we have $b_1 - a_1 + 1 = 2(b - a + 1)$. Thus, $|C_1| \leq 2(b_1 - a_1 + 1) = 4(b - a + 1)$. To complete the proof, we show that both C and $E_{[a,b]}$ are subsets of C_1 , which implies that $|V(S_{[a,b]})| \leq 2|E_{[a,b]} \cup C| \leq 2|C_1| \leq 8(b - a + 1)$. Clearly, $C \subseteq C_1$, as the sum in C goes through less summands than the sum defining C_1 . To show that $E_{[a,b]} \subseteq C_1$, consider $e \in E_{[a,b]}$. Suppose that $[a, b]$ is the left child of $\text{par}([a, b]) = [a_1, b_1]$. We show that $e \in \bigcup_{i=a_1}^{b_1-1} \Delta_i^-$. By Lemma 2, we have both $e \notin E_{[a_1, b_1]}$ and $e \notin E_{\text{right}([a_1, b_1])}$. Thus, the edge e is deleted in some version G_j for $j \in [b + 1, b_1]$, which means $e \in \Delta_{j-1}^-$. Analogously we prove that if $[a, b] = \text{right}([a_1, b_1])$ then $e \in \bigcup_{i=a_1+1}^{b_1} \Delta_i^+$. Hence $e \in C_1$. \square

To build T we use a recursive procedure $\text{COMPUTE-TREE}(a, b)$, which computes the subtree rooted at $T_{[a,b]}$. It produces each graph $S_{[a,b]}$ based on $S_{\text{par}([a,b])}$ by applying reduction and contraction. During the computation of T , we maintain an auxiliary table **repr**, fulfilling the following invariant: both at the beginning and at the end of the call $\text{COMPUTE-TREE}(a, b)$, **repr** $[v]$ is the vertex of $S_{\text{par}([a,b])}$ representing $v \in V$, if such vertex exists. Initially, we have **repr** $[v] = v$, which does not break the invariant, as we have previously set $\text{par}([1, t]) = [0, \infty]$ and $S_{[0, \infty]} = (V, \emptyset)$. The **repr** table is used implicitly by the procedures **REDUCE** and **CONTRACT** to map the endpoints of edges from $E_{[a,b]}$, Δ_i^+ and Δ_i^- to vertices of $S_{\text{par}([a,b])}$ in constant time.

All the computed tables use linear space and can be accessed in constant time, as we can identify the vertices of introduced graphs S_P with natural numbers $\{1, 2, \dots\}$ and the \perp value with 0. The total used space is asymptotically no more than the time spent on computing T , that is $O(m + t \log n)$.

```

1: procedure  $\text{COMPUTE-TREE}(a, b)$   $\triangleright [a, b]$  – elementary interval
2:    $P := \text{par}([a, b])$ 
3:   if  $b - a + 1 < n$  then  $\triangleright$  Reduction is only done for short elementary intervals.
4:      $C := \bigcup_{i=a+1}^b \Delta_i^+ \cup \bigcup_{i=a}^{b-1} \Delta_i^-$ 
5:      $U := \text{vertices of } V \text{ incident with any edge of } E_{[a,b]} \cup C$ 
6:      $(S', M') := \text{REDUCE}(S_P, E_{[a,b]} \cup C)$ 
7:   else
8:      $U := V$ 
9:      $(S', M') = (S_P, \text{id})$ 
10:  for  $u \in U$  do
11:     $\text{mem}[u] = \text{repr}[u]$   $\triangleright$  remember old repr values
12:  for  $u \in U$  do
13:     $\text{repr}[u] := M'(\text{repr}[u])$ 
14:   $(S'', M'') := \text{CONTRACT}(S', E_{[a,b]})$   $\triangleright M''$  maps  $V(S_P)$  to  $V(S_{[a,b]})$ .
15:   $S_{[a,b]} := S''$ 
16:  for  $k \in S_{[a,b]}$  do  $\triangleright$  initialize l and r pointers to  $\perp$ 
17:     $\text{l}_{[a,b]}[k] := \perp$ 
18:     $\text{r}_{[a,b]}[k] := \perp$ 
19:  for  $u \in U$  do
20:     $\text{repr}[u] := M''(\text{repr}[u])$ 
21:  for  $s \in S_P$  do  $\triangleright$  set the parent l and r pointers
22:     $s' := M'(s)$ 
23:    if  $s' \neq \perp$  then
24:       $s' := M''(s')$ 
25:    if  $[a, b] = \text{left}(P)$  then
26:       $\text{l}_P[s] := s'$ 
27:    else
28:       $\text{r}_P[s] := s'$ 
29:  if  $a < b$  then  $\triangleright$  compute the children
30:     $\text{mid} := \lfloor \frac{a+b}{2} \rfloor$ 
31:     $\text{COMPUTE-TREE}(a, \text{mid})$ 
32:     $\text{COMPUTE-TREE}(\text{mid} + 1, b)$ 
33:  for  $u \in U$  do  $\triangleright$  restore repr to the initial state
34:     $\text{repr}[u] := \text{mem}[u]$ 

```

Lemma 8. *The total running time of COMPUTE-TREE(1, t) is $O(m + t \log n)$.*

Proof. We first analyze the time spent in the call COMPUTE-TREE(a, b), excluding the work in recursive calls. Let C be $\bigcup_{i=a+1}^b \Delta_i^+ \cup \bigcup_{i=a}^{b-1} \Delta_i^-$. Thus $O(|C|) = O(b - a)$. Recall that the functions CONTRACT and REDUCE run in linear time. For $b - a + 1 \geq n$, we only perform contraction of $E_{[a,b]}$ in a graph of size $O(n)$, which requires $O(n + |E_{[a,b]}|)$ time. The amount of work for $b - a + 1 < n$ can be bounded by $O(|V(S_{\text{par}([a,b])})| + |C| + |E_{[a,b]}|)$, as REDUCE is passed the edges $C \cup E_{[a,b]}$.

To complete the proof, we sum these running times over all elementary intervals. The term $|E_{a,b}|$ appears in both cases and, by Lemma 4, we have $\sum_P E_P = O(m + t \log n)$, thus we can focus on the other summands. For the case $b - a + 1 \geq n$, the remaining work is $O(n)$, but there are only $O(\frac{t}{n})$ such intervals, so the total work is $O(t)$. On the other hand, if $b - a + 1 < n$, by Lemma 7, $O(|V(S_{\text{par}([a,b])})|) = O(b - a)$, so the total work is $O(b - a)$. Hence, the total work on each level of the tree such that its elementary intervals are shorter than n , is $O(t)$. The number of such levels is $O(\log n)$, which gives $O(t \log n)$ total time. The lemma follows. \square

Having computed T , the function COMP-ID allows us to access the component of some vertex v in G_c in time $O(\log t)$. However, as we now show, this can be speeded up to $O(\log n)$ time. Recall that $t = 2^B$. Let 2^D be the smallest power of 2 such that $2^D \geq n$ and fix some $k \in [0, 2^{B-D})$. Then, for each $c \in [k \cdot 2^D + 1, (k + 1) \cdot 2^D]$, the call $\text{COMP-ID}(1_{[0,\infty]}[v], 1, t, c)$ descends down T through the first $B - D$ levels in the same way, independent of c . We can thus add another preprocessing phase, building the table **shortcut**. For a vertex v and $0 \leq k < 2^{B-D}$, **shortcut** $[v][k]$ is defined to be a pair (s, P) such that for $c \in [k \cdot 2^D + 1, (k + 1) \cdot 2^D]$, $\text{COMP-ID}(1_{[0,\infty]}[v], 1, t, c)$, after going through at most $B - D$ levels of T , ends up in the interval P and $s \in V(S_P)$ represents v . There are only $O(t/n)$ allowed values of k , so the table **shortcut** has size $O(t)$.

The table can be computed by finding the components of each vertex v in all the graphs S_P from the first $B - D$ levels of the tree. As the component of v in S_P can be computed in constant time based on the component of v in $S_{\text{par}(P)}$, we spend $O(t/n)$ time for each v , and thus $O(t)$ time in total.

The optimized procedure COMP-ID starts by looking up the shortcut through first $B - D$ levels of T and then calls the original COMP-ID, starting at an elementary interval of length $O(n)$. Thus, its running time is $O(\log n)$.

3.2 2-edge-connectivity

As in the case of connectivity, we first show how to preprocess the graph in order to efficiently answer 2-edge-connectivity queries regarding individual versions. Our approach is similar to the idea of Section 3.1: we construct a data structure T containing graphs $S_{[a,b]}$, where $[a, b]$ is an elementary interval. Note that in the case of connectivity, the graphs $S_{[a,b]}$ do not contain any edges.

First, observe that contracting 2-edge-connected components yields a forest.

Lemma 9. *Let W be the set of 2-edge-connected components of some graph G . Define the graph $H = (W, F)$, where*

$$F = \{(w_1, w_2) : (u, v) \in E(G), u \in w_1, v \in w_2, w_1 \neq w_2\}.$$

Then, H is a forest.

Proof. Indeed, if there was a cycle $w_1 w_2 \dots w_k w_1$ in H , then the components w_1, \dots, w_k would form a single 2-edge-connected component. \square

In the case of 2-edge-connectivity, the graphs $S_{[a,b]}$ are forests of rooted trees, whose vertices represent some of the 2-edge-connected components of $G_{[a,b]}$. Each rooted tree in the forest represents a part of some (ordinary) connected component incident to some edges alive in the time interval $[a, b]$.

The vertices of $S_{[a,b]}$ are partitioned into two categories. A vertex $s \in V(S_{[a,b]})$ is a *simple vertex* if and only if it represents a single 2-edge-connected component of $G_{[a,b]}$. Otherwise, s is called a *path vertex* and it represents k ($k \geq 2$) 2-edge-connected components of $G_{[a,b]}$ — c_1, \dots, c_k — that form a “path”, i.e. for each $i \in [1, k)$ there is a single edge in G_P connecting some vertex of c_i and some vertex of c_{i+1} . We maintain the following invariants.

1. If s is a root of its tree, or its degree in S_P is other than 2, then it is a simple vertex.
2. A path vertex is never adjacent to another path vertex.

In particular, each path vertex is of degree 2 in $S_{[a,b]}$. Let s be a path vertex representing a path c_1, c_2, \dots, c_k of 2-edge-connected components of $G_{[a,b]}$. If s_1 is a parent of s and s_2 is a child of s , then the edge $(s_1, s) \in E(S_{[a,b]})$ is actually an edge between components s_1 and c_1 , while the edge $(s, s_2) \in E(S_{[a,b]})$ actually means (c_k, s_2) .

The components c_1, c_2, \dots, c_k of a path vertex $s \in V(S_{[a,b]})$ have the following property: for each $x \in [a, b]$, either c_1, c_2, \dots, c_k are actual 2-edge-connected components of G_x or they are all parts of a single larger 2-edge-connected component of G_x . Thus, the path vertex s allows us to trace components c_i in the descendants of $T_{[a,b]}$ in a uniform way. This in turn will later allow us to keep the graphs $S_{[a,b]}$ small.

Recall that in the case of connectivity, the call $\text{COMP-ID}(\mathbf{1}_{[0,\infty]}[v], 1, t, x)$ returns a pair (s, Q) ($s \in V(S_Q)$), where Q is the last interval on the path $[1, t] \rightarrow [x, x]$ in the segment tree such that v is represented with s in $V(S_Q)$. The introduction of path vertices forces us to modify what COMP-ID does, as we need to distinguish between distinct 2-edge-connected components that are represented by the same path vertex. Now for $x \in [a, b]$ and $k \in V(S_{[a,b]})$, we set $\text{COMP-ID}(k, a, b, x)$ to be the pair (s, Q) , where s is a representation of k in S_Q and Q is the last interval on the path $[a, b] \rightarrow [x, x]$ such that s is a simple vertex. To access the 2-edge-connected component of $v \in V$ in version G_x , we use the same call $\text{COMP-ID}(\mathbf{1}_{[0,\infty]}[v], 1, t, x)$.

The tables $\mathbf{l}_{[a,b]}$ and $\mathbf{r}_{[a,b]}$ are defined analogously: for $s \in V(S_{[a,b]})$, if $\mathbf{l}_{[a,b]}[s] = \perp$, then s is not represented in $S_{\text{left}([a,b])}$. Otherwise, $\mathbf{l}_{[a,b]}[s]$ is a vertex of $S_{\text{left}([a,b])}$ representing s . The table $\mathbf{r}_{[a,b]}$ is defined analogously.

Having defined graphs $S_{[a,b]}$, we now describe how they can be computed. As previously, we compute $S_{[a,b]}$ based on $S_{\text{par}([a,b])}$, by performing first the reduction and then the contraction. The reduction is again performed only if $b - a + 1 < n$.

The reduction proceeds in phases. The initial phases involve marking some nodes of $S_{\text{par}([a,b])}$, whereas the latter phases reduce the graph’s size. The path vertices never get marked; they can be instead merged with other path and simple vertices, forming “longer” path vertices of $S_{[a,b]}$.

Let C be again $\bigcup_{i=a+1}^b \Delta_i^+ \cup \bigcup_{i=a}^{b-1} \Delta_i^-$. In the first phase we mark all the vertices of $S_{\text{par}([a,b])}$ incident to edges in $E_{[a,b]} \cup C$. It might be the case that for some original edge (u, v) , u and v are already represented by the same s in $S_{\text{par}([a,b])}$ — then we just skip this edge. As a result, no more than $2|E_{[a,b]} \cup C|$ vertices of $S_{\text{par}([a,b])}$ are marked.

In the second phase we mark all the lowest common ancestors of marked vertices, that is, the vertices s such that in the first phase, the vertices from at least two distinct subtrees rooted at children of s , were marked. The common ancestors can be marked in linear time, using post-order traversal — we only need to store for each vertex s , whether any element of the subtree rooted at s was marked in the first phase. Additionally, we mark the root of every tree with at least one marked vertex.

Let us count the vertices marked after the second phase. Remove the subtrees with no marked vertices. Let q be the number of marked vertices of degree 2. If we replace all the degree 2 vertices with edges, we obtain a forest, where every vertex that is neither the leaf nor the root, has at least 2 children. Denote by l the number of leaves in this forest. Clearly, it has at most $2l$ vertices. However, every leaf could be marked only in the first phase and hence $l + q \leq 2|E_{[a,b]} \cup C|$, so $2l + q \leq 4|E_{[a,b]} \cup C|$.

Corollary 1. *After the second phase of the reduction, at most $4|E_{[a,b]} \cup C|$ vertices of $S_{\text{par}([a,b])}$ are marked.*

The third reduction phase removes the subtrees with no previously marked vertices. All the 2-edge-components represented by vertices from those components look exactly the same in $G_{[a,b]}$ as well as in all the individual versions G_a, \dots, G_b and thus need not be tracked in the descendants of $T_{[a,b]}$.

In the last phase we replace every remaining path of degree 2 unmarked vertices with a single path vertex. These vertices may include both simple and path vertices. However, neither of them has been marked, so for each $x \in [a, b]$, the underlying path of 2-edge-connected components c_1, \dots, c_g either remains unaltered in G_x or is a part of a single, larger 2-edge-component in G_x .

Since the number of such paths does not exceed the number of vertices marked so far, we end up with a forest S' of at most $8|E_{[a,b]} \cup C|$ vertices.

Each phase of the reduction can be implemented as a simple graph search, so the reduction takes time $O(|V(S_{\text{par}([a,b])})| + |E_{[a,b]}| + |C|)$.

After the reduction comes the contraction. We extend the forest S' with the edges $E_{[a,b]}$ alive in each of G_a, \dots, G_b . We merge the 2-edge-connected components found in this graph into new, simple vertices, obtaining a new graph S'' , which is again a forest. It may happen that some vertices of S' have not been merged into larger components in S'' . Every such vertex $s \in V(S')$ is a path vertex in S'' iff it is a path vertex in S' . The roots of trees of S'' are chosen arbitrarily, but keeping in mind that the trees should not be rooted at path vertices. The properly rooted S'' forms our graph S_P . Contraction can be implemented to work in time $O(|V(S_{\text{par}([a,b])})| + |E_{[a,b]}|)$.

Let us bound the time needed to compute $S_{[a,b]}$. If $b - a + 1 \geq n$, then the reduction is skipped and thus the time spent on building $S_{[a,b]}$ is $O(n + |E_{[a,b]}|)$. Otherwise, the reduction is performed and we spend $O(|V(S_{\text{par}([a,b])})| + |C| + |E_{[a,b]}|)$ time.

The asymptotic running time of building $S_{[a,b]}$ turns out to be exactly the same as in the case of connectivity. Thus, building a data structure T for representing 2-edge-connectivity takes the same time.

Corollary 2. *We can build a data structure T representing 2-edge-connectivity in a graph timeline in $O(m + t \log n)$ time. The space usage is $O(t \log n)$.*

The optimization allowing the evaluation of COMP-ID in time $O(\log n)$ applies here as well.

4 Answering forall queries

In this section we show how to extend the data structure T , so that it can be used for answering **forall** queries. The preprocessing for **forall** queries constitutes another phase, that we apply only after we computed the data structure T .

Let us begin with a simple observation. Assume that we want to answer a **forall**(u, w, a, b) query, where $[a, b]$ is an elementary interval. Then, if the same vertex of $S_{[a,b]}$ represents both u and w , then there is actually a path between u and w in $G_{[a,b]}$ and we can immediately give

a positive answer. However, the reverse relation is not true. It may happen that u and w are represented by distinct vertices in $S_{[a,b]}$, but are connected in each of G_a, \dots, G_b . Thus, our first goal in this section is to compute, for each two vertices in each of $S_{[a,b]}$, whether the vertices represented by them are connected in each of G_a, \dots, G_b .

For an elementary interval $[a, b]$, let $c_{[a,b]}(s, x)$, where $s \in V(S_{[a,b]})$, $x \in [a, b]$, be the result of the call $\text{COMP-ID}(s, a, b, x)$. Our goal is to compute for each vertex $s \in S_{[a,b]}$ a *fingerprint*, that is, an integer $H_{[a,b]}(s) \in [1, |V(S_{[a,b]})|]$ with the following property: the sequences $c_{[a,b]}(s, a)c_{[a,b]}(s, a+1) \dots c_{[a,b]}(s, b)$ and $c_{[a,b]}(s', a)c_{[a,b]}(s', a+1) \dots c_{[a,b]}(s', b)$ are equal iff $H_{[a,b]}(s) = H_{[a,b]}(s')$.

To answer a $\text{forall}(u, v, a, b)$ query, where $[a, b]$ is an elementary interval, we first map u and v into vertices u' and v' of $S_{[a,b]}$ and then report a positive answer iff $H_{[a,b]}(u') = H_{[a,b]}(v')$. In order to handle arbitrary intervals, we decompose the query interval into $O(\log t)$ elementary intervals. The decomposition as well as the mapping can be implemented as a function $\text{FORALL-AUX}(s_1, s_2, x, y, a, b)$, whose pseudocode is given in Appendix A. To answer a $\text{forall}(u, v, x, y)$ query we execute $\text{FORALL-AUX}(\mathbf{1}_{[0,\infty]}[u], \mathbf{1}_{[0,\infty]}[w], x, y, 1, t)$.

Let us now describe the computation of fingerprints. They are computed in a bottom-up fashion, starting from the leaves of T .

Lemma 10. *Let $P = [a, b]$ be an elementary interval and $s \in V(S_P)$. Define:*

$$\tilde{H}_P(s) = \begin{cases} (s, 0) & \text{if } \mathbf{1}_P(s) = \perp \text{ or } \mathbf{r}_P(s) = \perp \\ (H_{\text{left}(P)}(\mathbf{1}_P[s]), H_{\text{right}(P)}(\mathbf{r}_P[s])) & \text{otherwise.} \end{cases}$$

Then $c_P(s_1, a) \dots c_P(s_1, b) = c_P(s_2, a) \dots c_P(s_2, b)$ iff $\tilde{H}_P(s_1) = \tilde{H}_P(s_2)$.

Proof. If $\mathbf{1}_P[s_1] = \perp$, then $c_P(s_1, a) = (s_1, P)$ and for each $s_2 \in V(S_P)$ such that $s_1 \neq s_2$, we have $c_P(s_2, a) \neq (s_1, P)$. The pair $\tilde{H}_P(s_1) = (s_1, 0)$ is unique among the pairs $\tilde{H}_P(s)$, so we have $c_P(s_1, a) \dots c_P(s_1, b) = c_P(s_2, a) \dots c_P(s_2, b)$ iff $s_1 = s_2$. Analogously, if $\mathbf{r}_P[s_1] = \perp$, then $c_P(s_1, b) \neq c_P(s_2, b)$ for $s_1 \neq s_2$ and thus s_1 is given a unique pair $\tilde{H}_P(s_1) = (s_1, 0)$. Therefore, if $\mathbf{1}_P[s_1] = \perp$ or $\mathbf{r}_P[s_1] = \perp$, then $\tilde{H}_P(s_1) = \tilde{H}_P(s_2)$ is equivalent to $s_1 = s_2$.

It remains to consider the case, when s_1 and s_2 are represented in both $S_{\text{left}(P)}$ and $S_{\text{right}(P)}$. Let $m = \lfloor (a+b)/2 \rfloor$. By the definition of the pair $\tilde{H}_P(s)$ we have $c_P(s_i, a) \dots c_P(s_i, m) = c_{\text{left}(P)}(\mathbf{1}_P[s_i], a) \dots c_{\text{left}(P)}(\mathbf{1}_P[s_i], m)$ as well as $c_P(s_i, m+1) \dots c_P(s_i, b) = c_{\text{right}(P)}(\mathbf{r}_P[s_i], m+1) \dots c_{\text{right}(P)}(\mathbf{r}_P[s_i], b)$. The sequences $c_P(s_1, a) \dots c_P(s_1, b)$ and $c_P(s_2, a) \dots c_P(s_2, b)$, are equal exactly when their corresponding halves are equal, that is, by the definition of fingerprints H , iff $H_{\text{left}(P)}(\mathbf{1}_P[s_1]) = H_{\text{left}(P)}(\mathbf{1}_P[s_2])$ and $H_{\text{right}(P)}(\mathbf{r}_P[s_1]) = H_{\text{right}(P)}(\mathbf{r}_P[s_2])$. \square

Observe that the pairs $\tilde{H}_P(s)$ from the above lemma satisfy the desired properties of fingerprints, with the exception that they are pairs of integers, not integers. Thus, in order to compute the values $H_P(s)$, it suffices to map the values of $\tilde{H}_P(s)$ into distinct positive integers (two pairs are assigned the same integer iff they are equal). As both numbers in each pair $\tilde{H}_P(s)$ are at most $O(|V(S_P)|)$ we may compute the mapping in linear time by using radix-sort algorithm. Note that this resembles the Karp-Miller-Rosenberg [6] algorithm. The total additional time and space used is $O(\sum_P |V(S_P)|) = O(t \log n)$. Thus, we obtain an $\langle O(m + t \log n), O(\log t) \rangle$ data structure for answering forall queries.

However, the query time can be made independent of the length of the timeline and speeded up to $O(\log n)$. In order to do that, we employ a shortcutting technique similar to the one used for finding connected components of vertices in individual graphs combined with a data structure for comparing the subwords of a given word.

Assume again that D is the smallest integer such that $2^D \geq n$. Observe that the **shortcut** table from Section 3 allows us to speed up $\text{forall}(u, v, x, y)$, where $[x, y] \subseteq [k \cdot 2^D + 1, (k+1) \cdot 2^D]$

for some $k \in [0, 2^{B-D})$. Let $\text{shortcut}[u][k] = (s_u, P_u)$ and $\text{shortcut}[v][k] = (s_v, P_v)$. We only need to perform the following steps:

1. If $P_u \neq P_v$, then the answer is **false**.
2. If $P_u = P_v$, but the length of P_u is more than 2^D , then the answer is **true** if and only if $s_u = s_v$.
3. Otherwise, $P_u = P_v = [k \cdot 2^D + 1, (k+1) \cdot 2^D]$ and the answer can be obtained by calling $\text{FORALL-AUX}(s_u, s_v, x, y, k \cdot 2^D + 1, (k+1) \cdot 2^D)$.

Only the third steps takes superconstant time, namely $O(\log 2^D) = O(\log n)$.

Consider the general query $\text{forall}(u, v, x, y)$, where $[x, y] \not\subseteq [k \cdot 2^D + 1, (k+1) \cdot 2^D]$, for any k . Let l_1 be the smallest integer such that $x < l_1 \cdot 2^D + 1$ and l_2 be the largest integer such that $l_2 \cdot 2^D < y$. Then, our query can be split into the conjunction of three queries: $\text{forall}(u, v, x, l_1 \cdot 2^D)$, $\text{forall}(u, v, l_2 \cdot 2^D + 1, y)$ and $\text{forall}(u, v, l_1 \cdot 2^D + 1, l_2 \cdot 2^D)$ (we assume the last query to be **true** if $l_1 = l_2$). The first two can be answered in $O(\log n)$ time, as discussed above. We deal with the third one in a different way. Assume $l_1 < l_2$. For any l , let $\text{shortcut}[u][l] = (s_u, P_u)$ and $\text{shortcut}[v][l] = (s_v, P_v)$. Define $h_u^l := (H_{P_u}(s_u), P_u)$. Observe that the answer to $\text{forall}(u, v, l \cdot 2^D + 1, (l+1) \cdot 2^D)$ is affirmative exactly iff $h_u^l = h_v^l$, which follows immediately from the definition of fingerprints. Thus, to answer $\text{forall}(u, v, l_1 \cdot 2^D + 1, l_2 \cdot 2^D)$, we need to check if the sequences $h_u^{l_1} h_u^{l_1+1} \dots h_u^{l_2-1}$ and $h_v^{l_1} h_v^{l_1+1} \dots h_v^{l_2-1}$ are equal. We use the following algorithm, described for instance in [3]. It uses the linear construction of a suffix array and the optimal range minimum query structure.

Lemma 11 ([3]). *There exists a data structure that, after a linear preprocessing of the word W , allows us to check in time $O(1)$ if two subwords of W are equal.*

Let $X_v = h_v^0 h_v^1 \dots h_v^{2^{B-D}-1}$ and let X be the concatenation $X_1 X_2 \dots X_n$. Notice that the length of X is $O(t)$. We build the data structure of Lemma 11 for the sequence X in $O(t)$ time. Hence, we can answer the query $\text{forall}(u, v, l_1 \cdot 2^D + 1, l_2 \cdot 2^D)$ by comparing the appropriate two subwords of X of length $l_2 - l_1$ in time $O(1)$.

Theorem 1. *There exists an $\langle O(m + t \log n), O(\log n) \rangle$ data structure for answering forall queries.*

4.1 2-edge-connectivity

As in the case of the connectivity relation, for each $s \in V(S_{[a,b]})$, we want to encode the entire history of what happens with s in each of the individual versions G_a, \dots, G_b . Since we introduced path vertices in the graphs $S_{[a,b]}$, the appropriate fingerprints need to be defined in a more subtle way.

We partition the vertices of a graph $S_{[a,b]}$ into three groups:

1. simple vertices,
2. *vanishing* path vertices. If s is a vanishing path vertex and it represents a path of 2-edge-connected components c_1, \dots, c_k , then for each $x \in [a, b]$, all the components c_1, \dots, c_k are parts of a single, larger 2-edge-connected component C_x in G_x ,
3. *non-vanishing* path vertices. $s \in V(S_{[a,b]})$ is a non-vanishing path vertex if there exists $x \in [a, b]$ such that the underlying 2-edge-connected components c_1, \dots, c_k are all actual 2-edge-connected components of G_x .

Let us assume that $P = [a, b]$ and $s \in V(S_P)$. Denote by $c_P(s, x)$ the result of $\text{COMP-ID}(s, a, b, x)$. As in the case of connectivity we define $H_P(s) \in \{1, \dots, |V(S_P)|\} \cup \{\perp\}$ to be the fingerprint of the sequence $c_P(s, a) \dots c_P(s, b)$. We set $H_P(s) = \perp$ only if s is a non-vanishing path vertex. Otherwise, if s is simple or a vanishing path vertex, $H_P(s)$ is an integer.

There is a reason why a non-vanishing path vertex is not assigned an integer fingerprint: if vertices $v, w \in V$, $v \neq w$ are both represented by a non-vanishing path vertex s in S_P , then the sequences $c_{[1,t]}(1_{[0,\infty]}[v], a) \dots c_{[1,t]}(1_{[0,\infty]}[v], b)$ and $c_{[1,t]}(1_{[0,\infty]}[w], a) \dots c_{[1,t]}(1_{[0,\infty]}[w], b)$ might be different.

In order to define the fingerprints $H_P(s)$ based on the fingerprints for the children intervals $\text{left}(P)$ and $\text{right}(P)$, we first define the initial fingerprints $\tilde{H}_P(s)$. The values $\tilde{H}_P(s)$ have the same properties as the fingerprints $H_P(s)$ defined above, except that if $\tilde{H}_P(s) \neq \perp$, then $\tilde{H}_P(s)$ is a pair of integers from the range $[1, |V(S_P)|]$. The initial fingerprints can be computed according to the following rules, which implicitly decide whether a path vertex s is vanishing or non-vanishing.

1. If s is a simple vertex:

- if $1_P[s] = \perp$ or $r_P[s] = \perp$, then $\tilde{H}_P(s) = (s, 0)$. The fingerprint has to be unique in this case.
- Let $s_l = 1_P[s]$ and $s_r = r_P[s]$. If $H_{\text{left}(P)}[s_l] = \perp$ or $H_{\text{right}(P)}[s_r] = \perp$, then $\tilde{H}_P(s) = (s, 0)$. It is a consequence of s representing exactly a single 2-edge-connected component of G_x , for some $x \in [a, b]$.
- Otherwise, $\tilde{H}_P(s) = (H_{\text{left}(P)}(s_l), H_{\text{right}(P)}(s_r))$.

2. If s is a path vertex:

- if $1_P[s] = \perp$ or $r_P[s] = \perp$, then $\tilde{H}_P(s) = \perp$,
- Let $s_l = 1_P[s]$ and $s_r = r_P[s]$. If $H_{\text{left}(P)}[s_l] = \perp$ or $H_{\text{right}(P)}[s_r] = \perp$, then $\tilde{H}_P(s) = \perp$.
- Otherwise, $\tilde{H}_P(s) = (H_{\text{left}(P)}(s_l), H_{\text{right}(P)}(s_r))$.

Again, we can use radix-sort to convert each value $\tilde{H}_P(s)$ (distinct from \perp) to an integer $H_P(s)$ in the range $[1, |V(S_P)|]$.

We now sketch how to answer the query $\text{forall12}(u, v, x, y)$ in time $O(\log t)$. As in Section 4, we reduce this problem to answering $O(\log t)$ queries with the time period being an elementary interval. Let us focus on a single elementary interval P . For $w \in \{u, v\}$, denote by Q_w the last interval on the path $[1, t] \rightarrow P$ such that vertex w is represented in S_{Q_w} by a vertex s_w . Also, we denote by Q'_w the last interval on path $[1, t] \rightarrow P$ such that vertex w is represented in $S_{Q'_w}$ by a *simple* vertex s'_w . Denote the quadruple (Q_w, s_w, Q'_w, s'_w) by $\phi_P(w)$.

We have the following cases:

1. $Q_u \neq Q_v$. The answer is clearly **false**.
2. $Q_u = Q_v$, $Q_u \neq P$. As s_u or s_v could potentially be path vertices, the answer is positive if and only if $(Q'_u, s'_u) = (Q'_v, s'_v)$.
3. $Q_u = Q_v = P$.
 - (a) $H_P(s_u) \neq \perp$ and $H_P(s_v) \neq \perp$. The answer is **true** iff $H_P(s_u) = H_P(s_v)$.
 - (b) $H_P(s_u) = \perp$ or $H_P(s_v) = \perp$. The answer is the same as the result of the comparison $(Q'_u, s'_u) = (Q'_v, s'_v)$.

Once we have the fingerprints, all the above checks can be performed in $O(1)$ time, therefore we can answer **forall2** queries in $O(\log t)$ time.

In order to optimize the query time to $O(\log n)$, we adapt the technique used to speed up **forall** queries for connectivity. Assume again that D is the smallest integer such that $2^D \geq n$. First we show that we can answer the query **forall2**(u, v, x, y), where $[x, y] \subseteq [k \cdot 2^D + 1, (k+1) \cdot 2^D]$ for some $k \in [0, 2^{B-D})$, in time $O(\log n)$. We precompute the values $\phi_P(w)$ for each $w \in V$ and an elementary interval P not longer than 2^D . As the number of elementary intervals not longer than 2^D is $O(t/n)$, we precompute $O(t)$ values in total. $\phi_P(w)$ can be computed based on $\phi_{\text{par}(P)}(w)$ in constant time, so we spend $O(t)$ time on precomputation.

Let $[x, y] \subseteq [k \cdot 2^D + 1, (k+1) \cdot 2^D]$ and let P_1, \dots, P_p be the partition of $[x, y]$ into elementary intervals. As each P_i is a descendant of $[k \cdot 2^D + 1, (k+1) \cdot 2^D]$ in the segment tree, the value $\phi_{[k \cdot 2^D + 1, (k+1) \cdot 2^D]}(w)$ can be used as a starting point to compute the values $\phi_{P_i}(w)$. We need to descend only $B - D$ levels down the tree to compute the values $\phi_{P_i}(w)$. Thus, the time needed to answer **forall2**(u, v, x, y) is $O(B - D) = O(\log n)$ in this case.

To handle the general query **forall2**(u, v, x, y), let l_1 be the smallest integer such that $x < l_1 \cdot 2^D + 1$ and l_2 be the largest integer for which $l_2 \cdot 2^D < y$ holds. Our query can be split into the conjunction of three queries: **forall2**($u, v, x, l_1 \cdot 2^D$), **forall2**($u, v, l_2 \cdot 2^D + 1, y$) and **forall2**($u, v, l_1 \cdot 2^D + 1, l_2 \cdot 2^D$) (we assume the last query to be **true** if $l_1 = l_2$). The first two can be answered in $O(\log n)$ time, as discussed above. In order to answer the last query, we need to rephrase the check **forall2**($u, v, l \cdot 2^D + 1, (l+1) \cdot 2^D$) in terms of symbol equality $h_u^l = h_v^l$. Indeed, for $P_l = [l \cdot 2^D + 1, (l+1) \cdot 2^D]$ and $\phi_{P_l}(w) = (Q_w, s_w, Q'_w, s'_w)$ we can set:

$$h_w^l = \begin{cases} (Q_w, Q'_w, s'_w) & \text{if } Q_w \neq P_l \\ (Q_w, H_{P_l}(s_w)) & \text{if } Q_w = P_l \text{ and } H_{P_l}(s_w) \neq \perp \\ (Q_w, Q'_w, s'_w) & \text{if } Q_w = P_l \text{ and } H_{P_l}(s_w) = \perp \end{cases}$$

It can be easily verified that $h_u^l = h_v^l$ if and only if the previously described checks for answering **forall2**(u, v, x, y) where $[x, y] = [l \cdot 2^D + 1, (l+1) \cdot 2^D]$, produce a positive answer.

We answer the query **forall2**($u, v, l_1 \cdot 2^D + 1, l_2 \cdot 2^D$) by checking if the words $h_u^{l_1} \dots h_u^{l_2-1}$ and $h_v^{l_1} \dots h_v^{l_2-1}$ are equal. The needed words are all subwords of the word $X_1 \dots X_n$, where $X_v = h_v^0 h_v^1 \dots h_v^{2^{B-D}-1}$. The word length is $O(t)$. By Lemma 11, after additional preprocessing in $O(t)$ time, we can answer the query **forall2**($u, v, l_1 \cdot 2^D + 1, l_2 \cdot 2^D$) in constant time.

Theorem 2. *There exists an $\langle O(m + t \log n), O(\log n) \rangle$ data structure for answering **forall2** queries.*

5 Improved lower and upper bounds for exists queries

In this section we focus on **exists** queries. We first give improved conditional lower bounds for answering these queries, and then show an algorithm, whose running time matches one of the new bounds. As shown in [9], the problem of multiplying two Boolean $n \times n$ matrices can be reduced to the problem of answering $\Theta(n^2)$ **exists** queries about a graph timeline G^t , where $t = \Theta(n^2)$. Denote by $O(n^{\omega'})$ the time required to perform $n \times n$ Boolean matrix multiplication (BMM). Thus, unless $\omega' = 2$, it is not possible to develop a data structure, which after almost linear preprocessing answers **exists** queries in polylogarithmic time. In this section we give several new lower bounds.

Throughout this section, we repeatedly use ϵ to denote an arbitrarily small, positive number. The exact value of ϵ may vary and depend on the context. We also denote by $\delta(\epsilon)$ some other small positive number, dependent on ϵ .

Let us recall the somewhat informal, yet important, partition of algorithms into *algebraic* and *combinatorial*. The combinatorial algorithms do not make use of the fact that the matrices are defined over a ring, i.e., they do not use subtraction. No $O(n^{3-\epsilon})$ combinatorial algorithm is known for BMM.

We show a connection between the **exists** data structure and algorithmic problems related to detecting triangles in graphs. In the *triangle detection* problem we are given a graph $G = (V, E)$, where $|E| = m$, and the goal is to find three vertices $a, b, c \in V$ such that $(a, b), (a, c), (b, c) \in E$. The best known algorithm for triangle detection was given by Alon et al. [1] and works in $O(m^{1.41})$ time. The best combinatorial algorithm is folklore and runs in $O(m\sqrt{m})$ time. The following relation between triangle detection and BMM was shown in [11]:

Lemma 12. *An $O(m^{1.5-\epsilon})$ combinatorial algorithm for triangle detection implies an $O(n^{3-\delta(\epsilon)})$ combinatorial algorithm for BMM.*

The related problem is *triangle listing*, where we are asked to find c triangles in a graph with m edges. Pătraşcu [10] proved the following lemma.

Lemma 13. *If one can list m triangles from a graph with m edges in $O(m^{4/3-\epsilon})$ time, then there exists an $O(n^{2-\delta(\epsilon)})$ algorithm for 3-SUM.*

We now show a relation between triangle listing and **exists** queries.

Lemma 14. *The problem of listing c triangles in a graph with m edges can be reduced to answering $O(m + c \log n)$ **exists** queries in a timeline G^t of length $t = O(m)$ and no permanent edges.*

Proof. Let H be the input graph, in which we are supposed to list triangles. Moreover, let $V(H) = \{v_1, \dots, v_n\}$. We build a timeline G^t of graphs on vertex set $V(H)$ by processing vertices v_1, \dots, v_n one by one. First, we add an empty graph to G^t . Then, for a vertex v_i , we append $2 \deg_H(v_i)$ new versions to G^t ($\deg_H(v)$ denotes the degree of vertex v in H), which we call a *block* of vertex v_i . Within each block, we first create $\deg_H(v_i)$ new versions, at each step adding one more edge incident to v_i . The edges are added in arbitrary order. Then, we create $\deg_H(v_i)$ more versions by removing the edges incident to v_i . Note that the last graph in every block is empty, and in the middle graph the vertex v_i has degree $\deg_H(v_i)$. Let the block of a vertex v_i start at G_{a_i} and end at G_{b_i} .

Observe that we obtain a timeline G^t , where $t = 4m + 1$, as each edge of H is added and removed exactly twice. For each edge $(v_i, v_j) \in E(H)$, $i < j$, we can test if there is a triangle (v_i, v_j, v_k) , where $j < k$, with a single query **exists** $(v_i, v_j, b_j + 1, t)$. Indeed, the answer to such a query is positive iff there exists v_k such that there is a path from v_i to v_j in $G_{a_k + \deg_H(v_k) - 1}$. The path, along with the edge (v_i, v_j) , forms a triangle.

Note that the query **exists** (v_i, v_j, a_p, b_q) , for $j < p \leq q$, tells us if there is any triangle (v_i, v_j, v_k) such that $k \in [p, q]$. Thus, we may use a divide-and-conquer approach for listing triangles, which is based on the following observation. If we are looking for triangles such that $k \in [p, q]$, a negative answer to an **exists** $(v_i, v_j, a_p, b_{(p+q)/2})$ query allows us to halve the search interval. Hence, we can find all l vertices v_k such that (v_i, v_j, v_k) is a triangle in time $O(l \log n)$. The detailed procedure REPORT-TRIANGLES is given in Appendix A. \square

By combining Lemmas 12, 13 and 14, we obtain the following.

Theorem 3. *Let Ψ be a problem of answering $\Theta(t)$ **exists** queries about an arbitrary graph timeline G^t with no permanent edges.*

- An $O(t^{1.4})$ algorithm for Ψ implies an $O(t^{1.4})$ algorithm for triangle finding.
- An $O(t^{1.5-\epsilon})$ combinatorial algorithm for Ψ implies an $O(n^{3-\delta(\epsilon)})$ combinatorial algorithm for BMM.
- An $O(t^{4/3-\epsilon})$ algorithm for Ψ implies an $O(n^{2-\delta(\epsilon)})$ algorithm for 3-SUM.

In addition, we show that an **exists** data structure with preprocessing/query time product of $O(t^{2-\epsilon})$ and queries substantially faster than $O(\sqrt{t})$ implies a faster BMM algorithm.

Lemma 15. *Suppose there exists an $\langle O(t^{2-q-\epsilon}), O(t^q) \rangle$ combinatorial data structure for answering **exists** queries, where $q \in [0, \frac{1}{2})$ is a parameter. Then there exists an $O(n^{3-\delta(\epsilon)})$ combinatorial algorithm for BMM.*

Proof. We use the assumed algorithm for answering **exists** queries to develop an $O(n^{3-\delta(\epsilon)})$ combinatorial algorithm for BMM. By Lemma 12, in order to obtain such an algorithm, it suffices to show an $O(m^{\frac{3}{2}-\delta(\epsilon)})$ algorithm for triangle detection. This in turn, by Lemma 14, can be reduced to answering $O(m)$ **exists** queries in a graph timeline G^t with no permanent edges, where $t = O(m)$. To complete the proof, we show how to answer these queries in $O(m^{\frac{3}{2}-\delta(\epsilon)})$ total time.

Split G^t into blocks of length t^α consisting of consecutive versions, where $\alpha \in [0, 1]$ is to be set later. Let G_a, \dots, G_b be one of the blocks. First, we compute in time $O(t)$ the edges alive during entire block and contract the components formed by those edges. Next, we mark the components subject to any updates in the interval $[a, b]$. There are $O(t^\alpha)$ marked components. All the queries **exists**(u, v, c, d) such that $[c, d] \subseteq [a, b]$ and concerning the unmarked components, can be answered in constant time — it suffices to check whether components of u and v are the same. For the remaining $O(t^\alpha)$ components we build the assumed data structures. The initialization takes time $O(t^{\alpha(2-q-\epsilon)} + t)$ per individual block, which gives $O(t^{1+\alpha(1-q-\epsilon)} + t^{2-\alpha})$ time for $O(t^{1-\alpha})$ blocks.

We can answer an **exists**(u, v, x, y) query by going through at most $O(t^{1-\alpha})$ blocks intersecting $[x, y]$ and querying each block structure once. Thus, the query overhead is $O(t^{1-\alpha} \cdot t^{\alpha q}) = O(t^{1-\alpha(1-q)})$.

By setting $\alpha = \frac{1}{2-2q-\epsilon}$, we have $\alpha \in (\frac{1}{2}, 1]$. The exponent of the initialization time becomes less than $\frac{3}{2}$, as we have:

$$1 + \alpha(1 - q - \epsilon) = 1 + \frac{1 - q - \epsilon}{2 - 2q - \epsilon} < 1 + \frac{1 - q - \epsilon}{2 - 2q - 2\epsilon} = \frac{3}{2},$$

$$2 - \alpha < 2 - \frac{1}{2} = \frac{3}{2},$$

whereas the exponent of time needed to answer $O(t)$ queries is

$$1 + 1 - \alpha(1 - q) = 2 - \frac{1 - q}{2 - 2q - \epsilon} < 2 - \frac{1 - q}{2 - 2q} = \frac{3}{2}.$$

The lemma follows. □

What is interesting, we can give a combinatorial data structure, whose running time matches the above lower bound.

Theorem 4. *For every $0 \leq \alpha < 1$ there exists an $\langle O(m + \min(nt, t^{2-\alpha})), O(t^\alpha) \rangle$ data structure for answering **exists** queries. It uses $O(\min(nt, t^{2-\alpha}))$ space.*

Proof. If $n = o(t^{1-\alpha})$, then $nt = o(t^{2-\alpha})$ and the $\langle O(m+nt), O(1) \rangle$ data structure for answering **exists** queries from [9] is sufficient to finish the proof.

Let us assume that $n = \Omega(t^{1-\alpha})$. We first build the tree-like structure T in $O(m + t \log n)$ time. Let D be the largest integer such that $2^D \leq t^\alpha$. We split the graph timeline into blocks of size 2^{B-D} , which is roughly $t^{1-\alpha}$. Denote the versions of the i -th block by G_{a_i}, \dots, G_{b_i} . Observe that $[a_i, b_i]$ is an elementary interval.

Denote by $l_i(v)$ the pair (s, P) , where P is the last interval on a path from the root $[1, t]$ to $[a_i, b_i]$ such that $v \in V$ is represented by s in S_P .

Consider answering the *block query exists*(u, v, x, y), where $[x, y] \subseteq [a_i, b_i]$. Set $l_i(v) = (s_v, P_v)$ and $l_i(w) = (s_w, P_w)$. If $P_v \neq P_w$, then the answer is clearly **false**. Otherwise, if $P_v \neq [a_i, b_i]$, then neither v nor w are represented in $S_{[a_i, b_i]}$ and thus the answer is the same as the result of a comparison $s_v = s_w$. In the last case, when $P_v = P_w = [a_i, b_i]$, the result can vary. However, we can use the $\langle O(m+nt), O(1) \rangle$ data structure for answering **exists** queries from [9]. For each i , we build this structure for a graph with vertices $V(S_{[a_i, b_i]})$ and the timeline induced by the subsequent edge updates in graphs G_{a_i}, \dots, G_{b_i} . The size of $V(S_{[a_i, b_i]})$ as well as the length of this timeline is $O(b_i - a_i) = O(t^{1-\alpha})$, so it can be initialized in time $O((t^{1-\alpha})^2)$ to answer queries in $O(1)$ time. In a single query, it gives us the desired answer to **exists**(s_v, s_w, x, y).

As a result, once we have found $l_i(v)$ and $l_i(w)$, we can answer the block query in constant time. Observe that a general **exists** query can be answered by issuing $O(t^\alpha)$ block queries. All the required $O(t^\alpha)$ $l_j(u)$ values can be computed by traversing the first D levels of T , which can be implemented to work in time $O(t^\alpha)$. We need to build $O(t^\alpha)$ **exists** data structures — each in time $O((t^{1-\alpha})^2)$, so the total initialization time is $O(m + t^{2-\alpha} + t \log n) = O(m + t^{2-\alpha})$. \square

6 Open problems

For **forall** and **forall2** queries, we gave an $\langle O(m + t \log n), O(\log n) \rangle$ data structure. What about the biconnectivity? Although it is possible to propose a similar tree-like structure that represents biconnectivity in individual versions, it seems hard to extend it to **forall2**-like queries. The main obstacle is biconnectivity relation on vertices not being an equivalence relation.

It would be also interesting to know whether even faster query (without sacrificing $O(t \log n)$ initialization time) is possible for **forall** queries.

Concerning **exists** queries, we proved that beating our trade-off structure in the domain of combinatorial algorithms implies a faster combinatorial matrix multiplication algorithm. However, is there a way to employ fast matrix multiplication to obtain a data structure for **exists** queries with preprocessing/query time product of $O(t^{2-\epsilon})$?

References

- [1] Noga Alon, Raphael Yuster, and Uri Zwick. Finding and counting given length cycles. *Algorithmica*, 17(3):209–223, 1997.
- [2] Arnaud Casteigts, Paola Flocchini, Walter Quattrociocchi, and Nicola Santoro. Time-varying graphs and dynamic networks. *IJPEDS*, 27(5):387–408, 2012.
- [3] Maxime Crochemore, Christophe Hancart, and Thierry Lecroq. *Algorithms on Strings*. Cambridge University Press, New York, NY, USA, 2007.

- [4] David Eppstein. Offline algorithms for dynamic minimum spanning tree problems. *J. Algorithms*, 17(2):237–250, 1994.
- [5] Jacob Holm, Kristian de Lichtenberg, and Mikkell Thorup. Poly-logarithmic deterministic fully-dynamic algorithms for connectivity, minimum spanning tree, 2-edge, and biconnectivity. *J. ACM*, 48(4):723–760, 2001.
- [6] Richard M. Karp, Raymond E. Miller, and Arnold L. Rosenberg. Rapid identification of repeated patterns in strings, trees and arrays. In *Proceedings of the 4th Annual ACM Symposium on Theory of Computing, May 1-3, 1972, Denver, Colorado, USA*, pages 125–136, 1972.
- [7] David Kempe, Jon M. Kleinberg, and Amit Kumar. Connectivity and inference problems for temporal networks. *J. Comput. Syst. Sci.*, 64(4):820–842, 2002.
- [8] Fabian Kuhn, Nancy A. Lynch, and Rotem Oshman. Distributed computation in dynamic networks. In Leonard J. Schulman, editor, *Proceedings of the 42nd ACM Symposium on Theory of Computing, STOC 2010, Cambridge, Massachusetts, USA, 5-8 June 2010*, pages 513–522. ACM, 2010.
- [9] Jakub Łącki and Piotr Sankowski. Reachability in graph timelines. In *Proceedings of the 4th Conference on Innovations in Theoretical Computer Science, ITCS '13*, pages 257–268, New York, NY, USA, 2013. ACM.
- [10] Mihai Pătraşcu. Towards polynomial lower bounds for dynamic problems. In *Proceedings of the Forty-second ACM Symposium on Theory of Computing, STOC '10*, pages 603–610, New York, NY, USA, 2010. ACM.
- [11] Virginia Vassilevska Williams and Ryan Williams. Subcubic equivalences between path, matrix and triangle problems. In *Proceedings of the 2010 IEEE 51st Annual Symposium on Foundations of Computer Science, FOCS '10*, pages 645–654, Washington, DC, USA, 2010. IEEE Computer Society.

A Omitted pseudocode

```

1: function COMP-ID( $w, a, b, c$ )  $\triangleright w \in V(S_{[a,b]}), c \in [a, b]$ 
2:   if  $a = b$  then
3:     return  $(w, [a, b])$ 
4:    $\text{mid} := \lfloor \frac{a+b}{2} \rfloor$ 
5:   if  $c \leq \text{mid}$  then
6:     if  $\mathbf{l}_{[a,b]}[w] = \perp$  then
7:       return  $(w, [a, b])$ 
8:     else
9:       return COMP-ID( $\mathbf{l}_{[a,b]}[w], a, \text{mid}, c$ )
10:  else
11:    if  $\mathbf{r}_{[a,b]}[w] = \perp$  then
12:      return  $(w, [a, b])$ 
13:    else
14:      return COMP-ID( $\mathbf{r}_{[a,b]}[w], \text{mid} + 1, b, c$ )

```

```

1: function FORALL-AUX( $s_1, s_2, x, y, a, b$ )
2:   if  $s_1 = s_2$  then
3:     return true
4:   if  $[x, y] = [a, b]$  then  $\triangleright [x, y]$  is elementary, we refer to fingerprints
5:     if  $H_{[a,b]}(s_1) = H_{[a,b]}(s_2)$  then
6:       return true
7:     else
8:       return false
9:    $\text{mid} := \lfloor \frac{a+b}{2} \rfloor$ 
10:  if  $x \leq \text{mid}$  then
11:    if  $\mathbf{l}_{[a,b]}[s_1] = \perp$  or  $\mathbf{l}_{[a,b]}[s_2] = \perp$  then
12:      return false
13:    if not FORALL-AUX( $\mathbf{l}_{[a,b]}[s_1], \mathbf{l}_{[a,b]}[s_2], x, \text{mid}, a, \text{mid}$ ) then
14:      return false
15:  if  $y > \text{mid}$  then
16:    if  $\mathbf{r}_{[a,b]}[s_1] = \perp$  or  $\mathbf{r}_{[a,b]}[s_2] = \perp$  then
17:      return false
18:    if not FORALL-AUX( $\mathbf{r}_{[a,b]}[s_1], \mathbf{r}_{[a,b]}[s_2], \text{mid} + 1, y, \text{mid} + 1, b$ ) then
19:      return false
20:  return true

```

```

1: procedure REPORT-TRIANGLES( $G$ )
2:   construct the timeline  $G^t$  from the proof of Lemma 14 along with numbers  $a_i, b_i$ 
3:   for  $(u, v) \in E(G)$  do  $\triangleright u < v$ 
4:     REPORT-INTERNAL( $u, v, v + 1, |V(G)|$ )
5:
6: procedure REPORT-INTERNAL( $u, v, x, y$ )
7:   if  $x > y$  then
8:     return
9:   if not exists( $u, v, a_x, b_y$ ) then
10:    return
11:  if  $x = y$  then
12:    output triangle  $(u, v, x)$ 
13:    if  $k$  triangles have been reported so far, stop
14:  return
15:   $\text{mid} := \lfloor \frac{x+y}{2} \rfloor$ 
16:  REPORT-INTERNAL( $u, v, x, \text{mid}$ )
17:  REPORT-INTERNAL( $u, v, \text{mid} + 1, y$ )

```
